

Question 3

We would like you to setup an environment on AWS with the following components:

Part 1 - An EC2 instance

When creating an EC2 instance,

Resource: <https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/concepts.html>

Instance type: t2.micro - since it was free tier

AMI: Linux

Created a key-pair and saved it as a .pem file.

Q. A script should start on boot that will read the ID of the instance or assign an ID if it is the first boot, the ID should be a 6-character hexadecimal.

Q. Register with a Lambda function to confirm ID and IP address, with a flag denoting if this is a new registration or an update of an old ID.

Part 2 - The Registration Function

Q. A Lambda function that acts as an API that takes an ID, IP Address and a flag (to signify a new registration or update)

What is a Lambda function?

Resource: <https://docs.aws.amazon.com/lambda/latest/dg/welcome.html>

(What is AWS Lambda? AWS Documentation)

A Lambda handler function is the entry point for AWS Lambda to start executing the code. There should be an event that triggers the Lambda function.

In this question, there are 2 events that trigger the Lambda functions,

- The registration function: The event trigger is the HTTP POST request made by the EC2 instance's startup script. The event includes the instance ID, IP address, and a flag indicating whether it is a new registration or an update.
- The key updating function

Creating a Lambda function using Python.

Resource: <https://docs.aws.amazon.com/lambda/latest/dg/lambda-python.html>

(Building Lambda functions with Python AWS Documentation) Contains the basic code using boto3 and the definition of a Lambda handler function.

```
import json
```

```
import boto3

def lambda_handler(event, context):
```

Q. The script should then, on loop, write to a time series database, of your choosing, the timestamp and a 'secret key' hash.

What do we need to run SQL statements in a Lambda function?

Resource:

<https://boto3.amazonaws.com/v1/documentation/api/latest/reference/services/rds-data.html>

(RDSDataService Documentation) Contains the definitions of the methods client() and execute_statement().

Says that you need to import boto3. Has a sample code to create a service client as well.

```
rds = boto3.client('rds-data')
```

Additional resources about the RDS service,

<https://boto3.amazonaws.com/v1/documentation/api/latest/reference/core/boto3.html>

Has the definition of the boto3 source code,

<https://boto3.amazonaws.com/v1/documentation/api/latest/modules/boto3.html#client>

-client(service_name, region_name=None, api_version=None, use_ssl=True, verify=None, endpoint_url=None, aws_access_key_id=None, aws_secret_access_key=None, aws_session_token=None, config=None)

-service_name (string) – The name of a service, e.g. 's3' or 'ec2'.

To take the ID, IP Address and a flag, the following code is used,

```
instance_id = event['id']
ip_address = event['ip']
flag = event['flag']
```

Q. Function should check the ID exists

- **If it does and it is an update call, update the IP address of the ID in the RDS, then return a true flag.**
- **If it doesn't and it is a new registration call, add it to the RDS, then return a true flag**
- **Otherwise, return a false flag**

First check the flag

Store a query to be used in the execute_statement function().

```
if flag == 'new':
    query = f"INSERT INTO Instances (InstanceID, IPAddress) VALUES ('{instance_id}', '{ip_address}')
```

```

elif flag == 'update':
    query = f"UPDATE Instances SET IPAddress='{ip_address}' WHERE
InstanceID='{instance_id}'"
else:
    return {'status': 'error', 'message': 'Invalid flag'}

```

Then, Check if the ID exists. Only if the ID exists, execute the queries.

Resource for execute_statement() function:

https://boto3.amazonaws.com/v1/documentation/api/latest/reference/services/rds-data/client/execute_statement.html

The execute_statement() function needs the database name, the resourceArn and the secretArn. These were defined at the top of the code.

```

database = 'TimeSeriesDB'
resourceArn =
'arn:aws:timestream:us-east-2:010928184835:database/TimeSeriesDB'
secretArn =
'arn:aws:secretsmanager:us-east-2:010928184835:secret:RDSSecret-p9BRKu'

```

The Arns were taken from the **database information**

ARN (Amazon Resource Name) are unique identifiers used to specify particular AWS resources across all of AWS.

arn: Indicates that this is an Amazon Resource Name.

partition: The AWS partition (e.g., aws for AWS regions, aws-cn for China, aws-us-gov for AWS GovCloud).

service: The AWS service (e.g., s3, ec2).

region: The AWS region (e.g., us-east-1, eu-west-1).

account-id: The AWS account ID.

resource-type: The type of resource (e.g., instance, bucket).

resource-id: The unique identifier for the resource.

```

# Check if the ID exists in the database
try:
    rds.execute_statement(
        secretArn=secretArn,
        database=database,
        resourceArn=resourceArn,
        sql=query

```

```

    )
    return {'status': 'success'}
except Exception as e:
    return {'status': 'error', 'message': str(e)}

```

Part 3: The key updating function

Q. A separate Lambda function that will read the time series data from the EC2 instance, confirm the secret key with the one stored in the RDS

```

query = "SELECT InstanceID, SecretKey FROM Instances"
instances = rds.execute_statement(
    secretArn=secretArn,
    database=database,
    resourceArn=resourceArn,
    sql=query
) ['records']

```

A query is written to get the instanceID and SecretKey columns of the Instances database in the RDS.

The execute_statement() function is used.

Resource:

https://boto3.amazonaws.com/v1/documentation/api/latest/reference/services/rds-data/client/execute_statement.html

The response of the execute_statement() function contains a dictionary with 'records' and 'columnMetadata' parameters. Here we are accessing only the records.

Q. It will then generate a new secret key and access an API to update the secret key stored on the EC2 instance.

The secret key is a hexadecimal string.

So to generate a new hexadecimal string, use sha256 from the hashlib module in python. Hashlib.sha256 uses the old secret key to create a new one. Making sure that the new key is derived from the old one but is very different from the old one as well.

```
new_secret_key = hashlib.sha256(stored_secret_key.encode()).hexdigest()
```

Send the POST request to the correct url of the instance and with a payload dictionary.

This requires an API to be created which was not done.